

A4W4-unsup-learning_resubmit-UCB

May 28, 2024

```
[29]: from tensorflow.keras import layers
      from tensorflow import keras
      import tensorflow as tf
      import sklearn
      from sklearn.model_selection import train_test_split
      from ast import literal_eval
      import matplotlib.pyplot as plt
      import pandas as pd
      import numpy as np
      from sklearn.cluster import AgglomerativeClustering, KMeans
      import itertools
      import tensorflow_datasets as tfds
      from scipy.sparse import csr_matrix
      from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
      from IPython.display import Image
      from wordcloud import WordCloud, STOPWORDS
      import inspect
      from bayes_opt import BayesianOptimization
      from sklearn.metrics import   

      ↪ConfusionMatrixDisplay, confusion_matrix, accuracy_score
      from scipy.interpolate import griddata
      from bayes_opt.logger import JSONLogger
      from bayes_opt.event import Events
      from bayes_opt.util import load_logs
```

```
[157]: path = 'learn-ai-bbc'
      train = pd.read_csv(path+'/BBC News Train.csv')
      test = pd.read_csv(path+'/BBC News Test.csv')
```

```
[ ]: table=[]
      table.append(['Method', 'Train Accuracy', 'Test Accuracy'])
```

EXPLORATORY DATA ANALYSIS

First let us load dataframe.info to get the size of the dataframe

```
[31]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1490 entries, 0 to 1489
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   ArticleId   1490 non-null   int64
1   Text         1490 non-null   object
2   Category     1490 non-null   object
dtypes: int64(1), object(2)
memory usage: 35.1+ KB
```

Now let us check number of categories

```
[32]: train['Category'].nunique()
```

```
[32]: 5
```


Let us see the total duplicate titles and examine deduplicated dataset through histogram

```
[158]: total_duplicate_titles = sum(train["Text"].duplicated())
print(f"There are {total_duplicate_titles} duplicate titles.")
train = train[~train["Text"].duplicated()]
```

There are 50 duplicate titles.

```
[34]: plt.hist(train['Category'])
```

```
[34]: (array([335.,  0., 234.,  0.,  0., 266.,  0., 342.,  0., 263.]),
array([0. , 0.4, 0.8, 1.2, 1.6, 2. , 2.4, 2.8, 3.2, 3.6, 4. ]),
<BarContainer object of 10 artists>)
```



```
[35]: train["Text"].apply(lambda x: len(x.split())).describe()
```

```
[35]: count    1440.000000
      mean      385.215972
      std       212.167106
      min        90.000000
      25%       253.000000
      50%       337.000000
      75%       468.250000
      max      3345.000000
      Name: Text, dtype: float64
```

Now let us make a function to check permutations of label order to get the right clusters

```
[36]: ## PERMUTATION CHECK Function as used in week 2
      pd.set_option('future.no_silent_downcasting', True)
      def permute_label(y, label):

          k = list(train['Category'].unique())
          n = list(range(5))
          acc_cop = 0
          for i in itertools.permutations(n):
```

```
d = dict(zip(k,i))
l = list(label.astype(str).replace(d,inplace=False))
acc = sklearn.metrics.accuracy_score(l,yp )
if acc_cop< acc:
    acc_cop = acc
    final_labels_order = i
#print(f"accuracy : {acc_cop}")
#print(f"final_labels : {final_labels_order}")
return acc_cop,final_labels_order
```

WordCloud is a tool to see the frequencies of words. Lets try this first without stopwords which means common english words like 'him, she' etc. are not to be used.

```
[37]: wc = WordCloud(background_color = "white",max_words =300).generate(train.Text.  
      ↪to_string())  
plt.imshow(wc)  
plt.axis("off")  
plt.show()
```



```
[38]: wc = WordCloud(background_color = "white",max_words=
      ↪=300,stopwords=['s','u','new']+list(STOPWORDS)).generate(train.Text.
      ↪to_string())
      plt.imshow(wc)
      plt.axis("off")
      plt.show()
```



```

        model = AgglomerativeClustering(n_clusters = 5,linkage = linkage,
↪metric = affinity, )
        yp = model.fit(tfidf.toarray()).labels_
        ac, final_labels = permute_label(yp,train.Category)
        l.append([linkage,affinity,ac,final_labels,yp,model])
    except:
        pass
print(sorted(l,key=lambda x: x[2])[-1][:4])

```

```

linkage affinity accuracy final label
['ward', 'euclidean', 0.8756944444444444, (3, 4, 1, 0, 2)]

```

Final label is used to calculate the test result. We also append the test dataset to retrain the model.

```

[41]: tfidf = tfidf_vectorizer.fit_transform(pd.concat([train['Text'],test['Text']]))
l = (sorted(l,key=lambda x: x[2])[-1][-1]).fit_predict(tfidf.toarray()).
↪astype(str)

```

```

[42]: c=0
for i in final_labels:
    l[l==str(i)] = train.Category.unique()[c]
    c+=1

```

```

[43]: test['Category'] = l[1440:]
test[['ArticleId','Category']].to_csv('Solution_Agglomerative.csv',index=False)
!kaggle competitions submit -c learn-ai-bbc -f Solution_Agglomerative.csv -m
↪"Message"

```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /Users/shaurya/.kaggle/kaggle.json'

Warning: Looks like you're using an outdated API Version, please consider updating (server 1.6.14 / client 1.6.12)

100% | 9.13k/9.13k [00:01<00:00, 8.91kB/s]

Successfully submitted to BBC News Classification

TEST SET RESULT only with TRAIN DATA

```

[44]: Image('Desktop/Screenshot 2024-05-15 at 5.41.08 AM.png')

```

```

[44]:  Solution_Agglomerative.csv 0.20136 0.20136 ☐
Complete (after deadline) · 2h ago · Message

```


TEST SET RESULT with TEST DATA

```

[45]: Image("Desktop/Screenshot 2024-05-15 at 7.46.20 AM.png")

```

```

[45]:  Solution_Agglomerative.csv 0.48299 0.48299 ☐
Complete (after deadline) · 22s ago · Message


```

```
[46]: table.append(['Hierarchial Clustering',0.88,0.48])
```

The result on train set is good. The test set result is average.

```
[47]: ##MODEL KMEANS
tfidf_vectorizer = TfidfVectorizer(
    max_df=0.5,
    min_df=20,
    stop_words="english",
    max_features = 1000
)
tfidf = tfidf_vectorizer.fit_transform(train['Text'])
l = []
for i in range(100):
    model = KMeans(n_clusters = 5, random_state = i).fit(tfidf)
    label_model = model.labels_

    label = train['Category']
    acc, final_labels = permute_label(label_model,label)
    plt.scatter(i,acc)
    l.append([acc,final_labels,model])
plt.xlabel('random state')
plt.ylabel('accuracy')
plt.show()
acc,final_labels,model = sorted(l, key = lambda x:x[0])[-1]
print(f"accuracy",acc)
print(f"final_labels",final_labels)
```



```
accuracy 0.9354166666666667
final_labels (2, 3, 1, 4, 0)
```

We change the random state as the Kmeans method is giving different results every time. Then we store the best model. The result on the train set is good. The final label is used to calculate the test set result as above and gives us much better result than hierarchial clustering

```
[48]: tfidf = tfidf_vectorizer.fit_transform(pd.concat([train['Text'],test['Text']]))
l = model.predict(tfidf.toarray()).astype(str)
c=0
for i in final_labels:
    l[l==str(i)] = train.Category.unique()[c]
    c+=1
test['Category'] = l[1440:]
test[['ArticleId','Category']].to_csv('Solution_kMeans.csv',index=False)
!kaggle competitions submit -c learn-ai-bbc -f Solution_kMeans.csv -m "Message"
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /Users/shaurya/.kaggle/kaggle.json'

Warning: Looks like you're using an outdated API Version, please consider updating (server 1.6.14 / client 1.6.12)

```
100%|          | 8.96k/8.96k [00:01<00:00, 8.95kB/s]
```


Successfully submitted to BBC News Classification

TEST SET RESULT

```
[49]: Image( "Desktop/Screenshot 2024-05-15 at 8.07.10 AM.png")
```

[49]:

	Solution_kMeans.csv Complete (after deadline) · 1m ago · Message	0.55646	0.55646	<input type="checkbox"/>
---	--	----------------	----------------	--------------------------

Both the unsupervised learning methods are similar in performance but Kmeans beats Hierarchical Clustering

Summary of results

```
[50]: table.append(['KMeans',0.94,0.55])
```

```
[51]: pd.DataFrame(table[1:],columns=table[0][:]).set_index('Method')
```

```
[51]:
```

	Train Accuracy	Test Accuracy
Method		
Hierarchical Clustering	0.88	0.48
KMeans	0.94	0.55

MATRIX FACTORIZATION

We will use non negative matrix factorization to to get 5 components corresponding to 5 categories in the data. These components will be classified based on their weights for each document, i.e. the weight with the highest value for each document will be the category of the document. Both train and test data will be utilized after finding the correct label order, i.e., the one which matches the ground truth, through label_permute function.

```
[52]: def label_permute_2(yp,Categ):
    l =[]
    for i in itertools.permutations(train["Category"].unique()):
        lookup = tf.keras.layers.StringLookup(vocabulary =_
↪i,output_mode="one_hot",num_oov_indices = 0)
        label_mh = [lookup(label).numpy() for label in Categ]
        l.append([sklearn.metrics.accuracy_score(yp,label_mh),i])
    return sorted(l,key = lambda x : x[0])[-1]

def nmf_opt(tfidf_vectorizer, df1=None, Train=False):
    """
    Parameters:
    tfidf_vectorizer is the tfidf feature vector
    df1 is the test set when train = false
    and is none when train = True

    Return acc, label when train = true
```

```

matrix W when df1 = Test
"""
tfidf = tfidf_vectorizer.fit_transform(df1['Text'])
C = csr_matrix(tfidf)
X = sklearn.decomposition.NMF(n_components=5,solver='mu',init = 'nndsvda',
                             beta_loss="kullback-leibler",alpha_W=0.
↪00005,alpha_H=0.00005,
                             l1_ratio=0.5,random_state =4, max_iter =500).
↪fit(C)
W = (X.transform(C))
W = np.array([W[i]==max(W[i]) for i in range(len(df1))]) ##Since every word
↪vector is now divided into its weights for each category
#we decide to take the maximum strength for each sample and set that to 1
↪and rest to zero. Hence the strongest weight
#becomes the category.
if Train :
    out = label_permute_2(W,df1.Category.values)
else:
    out = W
return out

```

```


[53]: tfidf_vectorizer = TfidfVectorizer(
        max_df=0.5,
        min_df=20,
        stop_words="english",
        max_features=1000)
acc, label_order = nmf_opt(tfidf_vectorizer,df1 = train, Train = True)
print(f"label :{acc,label_order}")
result = nmf_opt(tfidf_vectorizer,df1 = train)
arr_res = np.array([result[i]==max(result[i]) for i in range(len(result))]).
↪astype(int)
yp = [np.take(label_order,np.argwhere(arr_res[i] == 1.0)[..., 0])[0] for i in
↪range(len(result))]
print(f"acc ; {accuracy_score(yp[:1440],list(train.Category))}")
ConfusionMatrixDisplay(confusion_matrix(yp[:1440],list(train.Category),labels =
↪label_order ),display_labels = label_order).plot()
plt.show()

```

```

label :(0.9520833333333333, ('business', 'politics', 'sport', 'entertainment',
'tech'))
acc ; 0.9520833333333333

```



So this is a fair enough result. Let us now try hyperparameter optimization for the tfidf word processing algorithm.

```
[71]: l = []

pbounds = {'max_df' : (0.1,1), 'min_df' : (1,100),}
## Find best parameters
def bb_function(max_df,min_df):
    tfidf_vectorizer = TfidfVectorizer(
        max_df=int(max_df*100)/100.,
        min_df=int(min_df),
        stop_words="english",
        max_features=1000
    )
    result = nmf_opt(tfidf_vectorizer,df1 = train)
    arr_res = np.array([result[i]==max(result[i]) for i in range(len(result))]).
    ↪astype(int)
    yp = [np.take(label_order,np.argwhere(arr_res[i] == 1.0)[..., 0])[0] for i in
    ↪in range(len(result))]
    acc = accuracy_score(yp[:1440],list(train.Category))
    return acc
```

```

optimizer = BayesianOptimization(f = bb_function,pbounds = pbounds,random_state=
    ↪= 1)
# logger = JSONLogger(path="./logs.log")
# optimizer.subscribe(Events.OPTIMIZATION_STEP, logger)

optimizer.maximize(
    init_points=100,
    n_iter=1000,
)
print('')

```

```

[73]: load_logs(optimizer, logs=["logs.log.json"]);
      dic = optimizer.max['params']
      print(int(dic['max_df']*100)/100.)

```

0.24
1100

```

[74]: dic = optimizer.max['params']

```


```

[76]: tfidf_vectorizer = TfidfVectorizer(
      max_df=dic['max_df'],
      min_df=int(dic['min_df']),
      stop_words="english",
      max_features=1000,
      )
result = nmf_opt(tfidf_vectorizer,df1 = train)
arr_res = np.array([result[i]==max(result[i]) for i in range(len(result))]).
    ↪astype(int)
yp = [np.take(label_order,np.argwhere(arr_res[i] == 1.0)[..., 0])[0] for i in
    ↪range(len(result))]

ConfusionMatrixDisplay(confusion_matrix(yp,list(train.Category),labels =
    ↪label_order ),display_labels = label_order).plot()
print(f"accuracy:{accuracy_score(yp,list(train.Category))}")
plt.show()

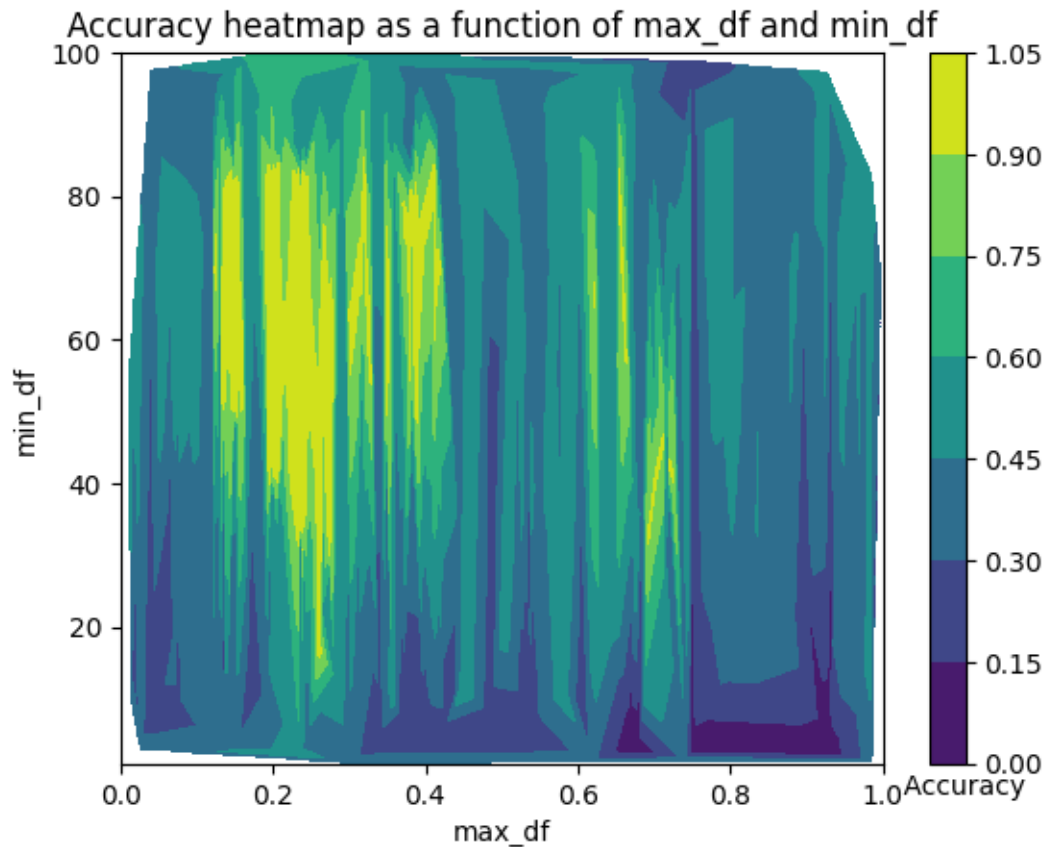
```

accuracy:0.9541666666666667



```
[77]: l=[]
a = ['target','params']
b = [[item[i]for item in optimizer.res]for i in a]
l = np.array([[k for j,k in i.items()] for i in b[1]])
```

```
[124]: x,y = np.linspace(0.1,1,1000),np.linspace(1,100,1000)
x,y = np.meshgrid(x,y)
g = griddata(points = (l[:,0],l[:,1]), values = b[0], xi = (x,y),method = 'linear')
plt.contourf(np.linspace(0,1,1000),np.linspace(1,100,1000),g.T)
plt.xlabel('max_df')
plt.ylabel('min_df')
plt.title('Accuracy heatmap as a function of max_df and min_df')
cb = plt.colorbar()
cb.ax.set_xlabel('Accuracy')
plt.show()
```



[]:

[]:

[]:

[]:

```
[130]: ## Check Test set
tfidf_vectorizer = TfidfVectorizer(
    max_df= int(dic['max_df']*100)/100,
    min_df=int(dic['min_df']),
    stop_words="english",
    max_features=1000)
result = nmf_opt(tfidf_vectorizer,df1 = pd.concat([train,test]))
arr_res = np.array([result[i]==max(result[i]) for i in range(len(result))]).
    ↳astype(int)
yp = [np.take(label_order,np.argmax(arr_res[i] == 1.0)[... , 0])[0] for i in_
    ↳range(len(result))]
test['Category'] = yp[1440:]
```

```
test[['ArticleId','Category']].to_csv('Solution_NMF.csv',index=False)
!kaggle competitions submit -c learn-ai-bbc -f Solution_NMF.csv -m "Message"
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /Users/shaurya/.kaggle/kaggle.json'
 Warning: Looks like you're using an outdated API Version, please consider updating (server 1.6.14 / client 1.6.12)
 100%| | 9.25k/9.25k [00:01<00:00, 8.71kB/s]
 Successfully submitted to BBC News Classification

TEST SET RESULT

[131]: Image("Desktop/Screenshot 2024-05-28 at 11.38.13 AM.png")

[131]:

	Solution_NMF.csv Complete (after deadline) · 20s ago · Message	0.93877	0.93877	<input type="checkbox"/>
---	--	----------------	----------------	--------------------------

[132]: table.append(['NMF',0.96,0.94])

[133]: pd.DataFrame(table[1:],columns=table[0][:]).set_index('Method')

[133]:

	Train Accuracy	Test Accuracy
Method		
Hierarchial Clustering	0.88	0.48
KMeans	0.94	0.55
NMF	0.96	0.94

I am happy with this score. To improve let us use another feature extraction method such as Word2Vec which uses word embedding to convert text into word vectors. From EDA we saw that max length is 3345

[134]:

```
word2vec = tf.keras.layers.TextVectorization(
    max_tokens=3345,
    standardize='lower_and_strip_punctuation',
    split='whitespace',
    output_mode='int',
    output_sequence_length=None,
    pad_to_max_tokens=False,
    encoding='utf-8',
)
```

[140]:

```
def label_permute_2(yp,Categ):
    l =[]
    for i in itertools.permutations(train["Category"].unique()):
        lookup = tf.keras.layers.StringLookup(vocabulary =_
        ↪i,output_mode="one_hot",num_oov_indices = 0)
```

```

        label_mh = [lookup(label).numpy() for label in Categ]
        l.append([sklearn.metrics.accuracy_score(yp,label_mh),i])
    return sorted(l,key = lambda x : x[0])[-1]

```

```

def nmf_opt(word2vec,df1=None,Train=False):
    ## NMF optimizer function

    """
    Parameters:
    tfidf_vectorizer is the tfidf feature vector
    df1 is the test set when train = false
    and is none when train = True

    Return acc, label when train = true
    matrix W when df1 = Test
    """

    word2vec.adapt(df1.Text)
    word2vec_df = word2vec(df1.Text)
    C = csr_matrix(word2vec_df)
    X = sklearn.decomposition.NMF(n_components=5,solver='mu',init = 'nndsvda',
                                beta_loss="kullback-leibler",alpha_W=0.
                                ↪0.00005,alpha_H=0.00005,
                                l1_ratio=0.5,random_state =4).fit(C)

    W = (X.transform(C))
    W = np.array([W[i]==max(W[i]) for i in range(len(df1))])
    if Train :
        out = label_permute_2(W,train['Category'].values)
    else:
        out = W
    return out

```

```

[179]: def remove_stopwords_word2vec_only(st):
        st = st.lower().split()
        m = [i for i in st if i not in STOPWORDS]
        return str(m)

    train_w2v = train.copy()
    train_w2v.Text = train.Text.apply(lambda x: remove_stopwords_word2vec_only(x))

```

```

[180]: ## Find best parameters
word2vec = tf.keras.layers.TextVectorization(
max_tokens=3345,
standardize='lower_and_strip_punctuation',
split='whitespace',
output_mode='int',
output_sequence_length=None,
pad_to_max_tokens=False,

```



```

encoding='utf-8',
)
acc, label_order = nmf_opt(word2vec,df1=train_w2v,Train=True)
result = nmf_opt(word2vec,df1 = train_w2v)
arr_res = np.array([result[i]==max(result[i]) for i in range(len(result))]).
    ↳astype(int)
yp = [np.take(label_order,np.argwhere(arr_res[i] == 1.0)[..., 0])[0] for i in
    ↳range(len(result))]

```

```


[181]: ConfusionMatrixDisplay(confusion_matrix(yp,list(train.Category),labels =
    ↳label_order ),display_labels = label_order).plot()

```

```

[181]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x3d5617020>

```



```

[182]: print(F"Train set -- accuracy, {acc}")
        print(F"Label order -- :{label_order}")

```

Train set -- accuracy, 0.30833333333333335

Label order -- :('entertainment', 'sport', 'business', 'politics', 'tech')

The accuracy is bad using this method. The label order is also wrong which we know from previous results. However, the mismatching is too bad to help.


SUPERVISED LEARNING


```
[189]: from sklearn.tree import DecisionTreeClassifier
       from sklearn.metrics import accuracy_score, recall_score
```

```
[212]: ##use KNN on the data
tfidf_vectorizer = TfidfVectorizer(
    max_df=0.5,
    min_df=20,
    stop_words="english",
    max_features = 1000
)
tfidf = tfidf_vectorizer.fit_transform(train['Text'])
tfidf_test = tfidf_vectorizer.fit_transform(test['Text'])
##Check best value of number of nearest neighbors
l=[]
m=[]
for i in range(1,50):

    l.append([accuracy_score(sklearn.neighbors.KNeighborsClassifier(
        n_neighbors = i).fit(tfidf,train.Category).predict(tfidf),train.
↪Category),i])
    m.append(accuracy_score(sklearn.neighbors.KNeighborsClassifier(
        n_neighbors = i).fit(tfidf,train.Category).predict(tfidf),train.
↪Category))
print("max acc, neighbors" , (sorted(l,key =lambda x: x[0]))[-1])
nn = (sorted(l,key =lambda x: x[0]))[-1][1]
##Train splitting accuracy check
plt.plot(range(1,50),m)
plt.xlabel('Nearest Neighbors')
plt.ylabel('Accuracy')
plt.title('Accuracy vs number of nearest neighbors')
y_pred = (sklearn.neighbors.KNeighborsClassifier(
    n_neighbors = nn).fit(tfidf,train.Category).predict(tfidf))
ConfusionMatrixDisplay(confusion_matrix(y_pred,list(train.Category))).plot()
plt.show()
for i in np.arange(1,11,2)*0.1:
    model = sklearn.neighbors.KNeighborsClassifier(
        n_neighbors = nn).fit(*sklearn.utils.shuffle(tfidf,train.
↪Category,random_state =1 ,n_samples=int(1440*i)))
    l = ([model.score(tfidf,train.Category),(i)])
    print(" acc, split" , l)
test['Category']= sklearn.neighbors.KNeighborsClassifier(n_neighbors = nn).
↪fit(tfidf,train.Category).predict(tfidf_test)
test[['ArticleId','Category']].to_csv('Solution_KNN.csv',index=False)
```

```
max acc, neighbors [1.0, 1]
```





```
acc, split [0.875, 0.1]
acc, split [0.9236111111111112, 0.30000000000000004]
acc, split [0.9451388888888889, 0.5]
acc, split [0.9659722222222222, 0.7000000000000001]
acc, split [0.9888888888888889, 0.9]
```

TEST SET RESULT

[266]: `Image("Desktop/Screenshot 2024-05-15 at 5.32.37 AM.png")`

[266]:



Solution_KNN.csv
Complete (after deadline) · 2h ago

0.34149

0.34149



[184]: `table.append(['KNN', 1.0, 0.34])`

50% data gives 95% accuracy on train set. The train set accuracy is 100. The test set result is little lesser to the 2 typical unsupervised learning methods' results and similar to NMF.


[214]: `##Use Decision Tree`
`from sklearn.tree import DecisionTreeClassifier`

```

tfidf_vectorizer = TfidfVectorizer(
    max_df=0.5,
    min_df=20,
    stop_words="english",
    max_features = 1000
)
tfidf = tfidf_vectorizer.fit_transform(train['Text'])
tfidf_test = tfidf_vectorizer.fit_transform(test['Text'])
def build_dt(X,y,max_depth=None, max_leaf_nodes = None, ccp=0):
    ↵
    ↵clf=DecisionTreeClassifier(random_state=0,max_depth=max_depth,max_leaf_nodes=max_leaf_nodes)
    ↵return clf.fit(X,y)
clf1 = build_dt(tfidf,train.Category)
print('acc = ' , accuracy_score(clf1.predict(tfidf),train.Category))
y_pred = (clf1.fit(tfidf,train.Category).predict(tfidf))
ConfusionMatrixDisplay(confusion_matrix(y_pred,list(train.Category))).plot()
plt.show()
yp = (clf1.predict(tfidf_test))
test['Category'] = yp
test[['ArticleId','Category']].to_csv(path+'/Solution_clf.csv',index=False)

```

acc = 1.0



```
[215]: ##Check accuracy by strength of train set splitting
for i in np.arange(1,11,2)*0.1:
    model = build_dt(*sklearn.utils.shuffle(tfidf,train.Category,random_state_u
    ↪=1 ,n_samples=int(1440*i)))

    l = ([model.score(tfidf,train.Category),(i)])
    print(" acc, split" , l)
```

```
acc, split [0.6555555555555556, 0.1]
acc, split [0.79375, 0.30000000000000004]
acc, split [0.8708333333333333, 0.5]
acc, split [0.9361111111111111, 0.7000000000000001]
acc, split [0.975, 0.9]
```

TEST SET RESULT

```
[267]: Image("Desktop/Screenshot 2024-05-15 at 4.07.43 AM.png")
```

[267]:

 **Solution_clf.csv** 0.29251 0.29251 ☐

Complete (after deadline) - now

KNN works better than Decision Tree which overfits and uses 90% data to get more than 95% accuracy. The result is a little poor compared to unsupervised learning on the test set. The train set is completely fitted with 100% accuracy

References: [https://scikit-learn.org/stable/auto_examples/applications/plot_topics_extraction_with_nmf_lda.](https://scikit-learn.org/stable/auto_examples/applications/plot_topics_extraction_with_nmf_lda.html)
[glr-auto-examples-applications-plot-topics-extraction-with-nmf-lda-py](https://scikit-learn.org/stable/auto_examples/applications/plot_topics_extraction_with_nmf_lda.html)

```
[185]: table.append(['Decision Tree',1.0,0.29])
```

```
[186]: pd.DataFrame(table[1:],columns=table[0][:]).set_index('Method')
```

```
[186]:
```

	Train Accuracy	Test Accuracy
Method		
Hierarchial Clustering	0.88	0.48
KMeans	0.94	0.55
NMF	0.96	0.94
KNN	1.00	0.34
Decision Tree	1.00	0.29

```
[ ]:
```